

The figure displays eight plots arranged in a 2x4 grid, showing the Expected coverage (Y-axis, 0.0 to 1.0) versus the Nominal credibility α (X-axis, 0.0 to 1.0). Each plot includes a solid line representing the model's performance and a dashed line representing the ideal performance (y=x). Shaded regions indicate confidence intervals.

Top Row: CNN models with TARP coverage — 3-D (Ω_m, σ_8, w_0)

- Top Left:** CNN auto native-TFDS (RealNVP), N seeds = 3. The model's performance is slightly below the ideal line.
- Top Second:** CNN auto+cross (MAF), N seeds = 3. The model's performance is slightly above the ideal line.
- Top Third:** CNN auto+cross (RealNVP), N seeds = 3. The model's performance is slightly above the ideal line.
- Top Right:** CNN auto+cross (RealNVP, std), N seeds = 3. The model's performance is slightly above the ideal line.

Bottom Row: CNN and L1 models

- Bottom Left:** CNN auto-only (MAF), N seeds = 3. The model's performance is slightly below the ideal line.
- Bottom Second:** CNN auto-only (RealNVP), N seeds = 3. The model's performance is slightly below the ideal line.
- Bottom Third:** L1 auto+cross, N seeds = 3. The model's performance is slightly above the ideal line.
- Bottom Right:** L1 auto-only, N seeds = 3. The model's performance is slightly below the ideal line.

