

Per-patch distributions over the fiducial population: l1+product $\rightarrow$ RealNVP (N=1000) vs CNN ResNet18 $\rightarrow$ RealNVP (N=9000)
bars = medians $\cdot$ CNN posteriors are tighter+more stable; l1 has larger per-patch scatter

