

joint PDF (counts) vs σ_8 — 16 levels, BNT basis (dequantized training cache)

